

Supplementary: Model-level diagnostic profile

FCC/C-RWR are primary readouts; remaining columns provide safety-gate, reviewer-dispersion, and operational context.

	FCC	C-RWR	Safety gap	Judge SD	Open100	Runtime	Turns	Workspace
GPT-5.6 Sol	0.824	0.753	0.072	0.140	8.3%	1.7%	22.3	2.3
GPT-5.4	0.786	0.680	0.106	0.157	5.8%	0.8%	23.1	1.2
DeepSeek-V4-Pro	0.783	0.621	0.162	0.162	5.0%	1.7%	24.4	10.6
Qwen3.8-Max	0.773	0.619	0.154	0.189	6.7%	2.5%	24.5	3.7
MiniMax M3	0.748	0.586	0.161	0.184	6.7%	18.3%	22.6	2.0
GLM-5	0.740	0.598	0.141	0.160	0.8%	1.7%	17.1	7.4
DeepSeek-V3.2	0.738	0.568	0.170	0.184	5.8%	0.8%	24.0	8.4
Kimi K2	0.694	0.510	0.184	0.181	2.5%	0.8%	19.6	3.5
Doubao-Seed-2.0-lite	0.651	0.471	0.181	0.184	1.7%	0.8%	17.3	1.8
Qwen3-30B-A3B-Instruct-2507	0.645	0.450	0.195	0.207	3.3%	2.5%	21.0	8.5

Safety gap = FCC minus C-RWR; larger values indicate a stronger clinical-error penalty. Judge SD is the mean within-trajectory FCC SD across four Judges.

Turn and workspace counts describe process behavior and are not quality scores. All models were evaluated on the same 120 cases.